

# Final Project: Bandit Learning

Ziyu Shao

School of Information Science and Technology  
ShanghaiTech University

May 16, 2023

# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# Example of Online Problem

- A senior undergraduate student of SIST in ShanghaiTech University have applied many top graduate programs around the world.
- He will receive several offers in a sequential order. After looking at an offer, he must decide whether to accept it (and terminate the process) or to reject it. Once rejected, an offer is lost.
- Suppose that the only information he has at any time is the relative rank of the present offer compared with previously ones.
- His objective is to maximize the probability of selecting the best offer.
- Please help him to find the optimal policy and compute the corresponding probability.

# Offline Problems vs. Online Problems

- Offline Problems
  - ▶ Assumed complete information: for example, know the whole input in advance
  - ▶ Example: the shortest path problem (topology & costs between nodes are given)
- Online Problems: decision, optimization and learning
  - ▶ Having no or incomplete knowledge of the future
  - ▶ Multiple decisions are made sequentially based on a piece-by-piece input
  - ▶ The sequential decisions are irrevocable
  - ▶ Example: playing the Chess/Go
- Bandit model & algorithms for online problems

# One-Armed Bandit

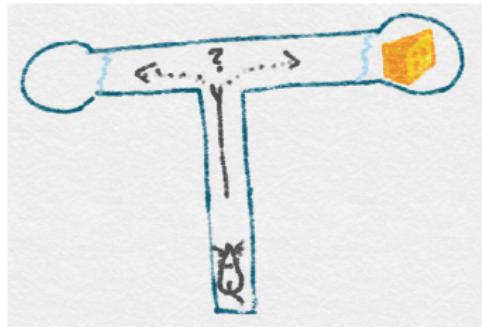


# Multi-Armed Bandit



# History

First bandit algorithm proposed by [Thompson \(1933\)](#)



[Bush and Mosteller \(1953\)](#) were interested in how mice behaved in a T-maze

# Widely Applications

- Clinical trials/dose discovery
- Recommendation systems (movies/news/etc)
- Advertisement placement
- A/B testing
- Monte Carlo tree search algorithm (game playing)
- Resource allocation
- Network routing
- Dynamic pricing (e.g.,for Amazon products)
- Ranking (e.g.,for search)

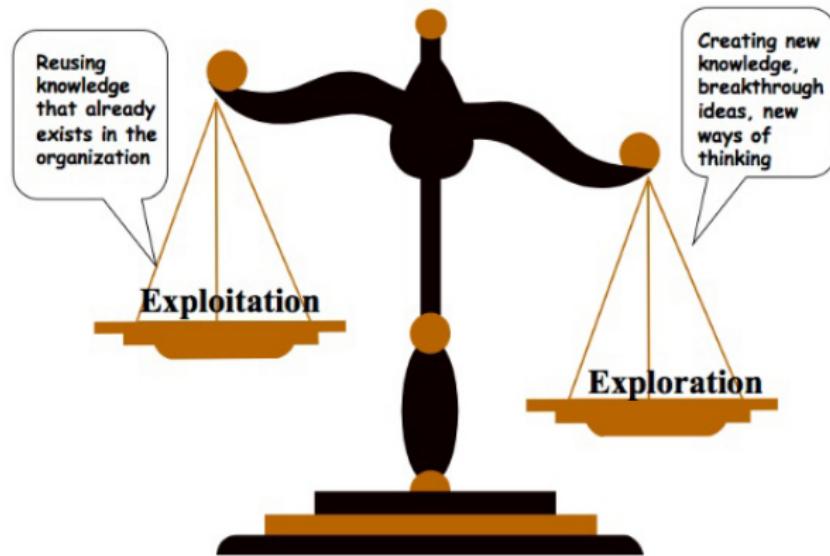
# Why Bandit?

- A perfect model for online decision making under uncertainty
- Isolate an important component of reinforcement learning: exploration-vs-exploitation
- Rich and beautiful mathematically

# Exploration vs. Exploitation Dilemma

- Online decision-making involves a fundamental choice
  - ▶ **Exploitation:** Make the best decision given current information
  - ▶ **Exploration:** Gather more information
- The best long-term strategy may involve short-term sacrifices
- Gather enough information to make the best overall decisions

# Exploration vs. Exploitation Dilemma



# Examples in Real Life

- Restaurant Selection
  - ▶ **Exploitation:** Go to your favorite restaurant
  - ▶ **Exploration:** Try a new restaurant

# Example of A Two-Armed Bandit



Time	1	2	3	4	5	6	7	8	9	10	11	12
Left arm	\$1	\$0			\$1	\$1	\$0					
Right arm				\$1	\$0							

# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# Basic Setting

- We consider a time-slotted bandit system ( $t = 0, 1, 2, \dots$ ) with three arms.
- We denote the arm set as  $\{1, 2, 3\}$ . Pulling each arm  $j$  ( $j \in \{1, 2, 3\}$ ) will obtain a reward  $r_j$ , which satisfies a Bernoulli distribution with mean  $\theta_j$  ( $\text{Bern}(\theta_j)$ ), i.e.,

$$r_j = \begin{cases} 1, & \text{w.p. } \theta_j, \\ 0, & \text{w.p. } 1 - \theta_j, \end{cases}$$

- $\theta_j$  are parameters within  $(0, 1)$  for  $j \in \{1, 2, 3\}$ .

# Basic Setting

- Now we run this bandit system for  $N$  ( $N \gg 3$ ) time slots.
- At each time slot  $t$ , we choose one and only one arm from these three arms, which we denote as  $I(t) \in \{1, 2, 3\}$ . Then we pull the arm  $I(t)$  and obtain a reward  $r_{I(t)}$ .
- Our objective is to find an optimal policy to choose an arm  $I(t)$  at each time slot  $t$  such that the expectation of the aggregated reward is maximized, i.e.,

$$\max_{I(t), t=1, \dots, N} \mathbb{E} \left[ \sum_{t=1}^N r_{I(t)} \right].$$

# Basic Setting

- If we know the values of  $\theta_j, j \in \{1, 2, 3\}$ , this problem is trivial.
- Since  $r_{I(t)} \sim \text{Bern}(\theta_{I(t)})$ ,

$$\mathbb{E} \left[ \sum_{t=1}^N r_{I(t)} \right] = \sum_{t=1}^N \mathbb{E}[r_{I(t)}] = \sum_{t=1}^N \theta_{I(t)}.$$

- Let  $I(t) = I^* = \arg \max_{j \in \{1, 2, 3\}} \theta_j$  for  $t = 1, 2, \dots, N$ , then

$$\max_{I(t), t=1, \dots, N} \mathbb{E} \left[ \sum_{t=1}^N r_{I(t)} \right] = N \cdot \theta_{I^*}.$$

- However, in reality, we do not know the values of  $\theta_j, j \in \{1, 2, 3\}$ . We need to estimate the values  $\theta_j$  via empirical samples, and then make the decisions at each time slot.

# Classical Bandit Algorithms

- $\epsilon$ -Greedy
- Upper Confidence Bound (UCB)
- Thompson Sampling

# $\epsilon$ -Greedy Algorithm

---

## Algorithm 1 $\epsilon$ -greedy Algorithm

---

**Initialize**  $\hat{\theta}(j) = 0$ ,  $\text{count}(j) = 0, j \in \{1, 2, 3\}$

1: **for**  $t = 1, 2, \dots, N$  **do**

2:

$$I(t) \leftarrow \begin{cases} \arg \max_{j \in \{1, 2, 3\}} \hat{\theta}(j) & w.p. 1 - \epsilon \\ \text{randomly chosen from } \{1, 2, 3\} & w.p. \epsilon \end{cases}$$

3:  $\text{count}(I(t)) \leftarrow \text{count}(I(t)) + 1$

4:  $\hat{\theta}(I(t)) \leftarrow \hat{\theta}(I(t)) + \frac{1}{\text{count}(I(t))} [r_{I(t)} - \hat{\theta}(I(t))]$

5: **end for**

---

# UCB Algorithm

---

## Algorithm 2 UCB Algorithm

---

```
1: for  $t = 1, 2, 3$  do
2:    $I(t) \leftarrow t$ 
3:    $\text{count}(I(t)) \leftarrow 1$ 
4:    $\hat{\theta}(I(t)) \leftarrow r_{I(t)}$ 
5: end for
6: for  $t = 4, \dots, N$  do
7:
8:    $I(t) \leftarrow \arg \max_{j \in \{1, 2, 3\}} \left( \hat{\theta}(j) + c \cdot \sqrt{\frac{2 \log t}{\text{count}(j)}} \right)$ 
9:    $\text{count}(I(t)) \leftarrow \text{count}(I(t)) + 1$ 
10:   $\hat{\theta}(I(t)) \leftarrow \hat{\theta}(I(t)) + \frac{1}{\text{count}(I(t))} [r_{I(t)} - \hat{\theta}(I(t))]$ 
11: end for
```

---

# Thompson Sampling Algorithm

---

**Algorithm 3** Thompson sampling Algorithm

---

**Initialize** Beta parameter  $(\alpha_j, \beta_j), j \in \{1, 2, 3\}$

1: **for**  $t = 1, 2, \dots, N$  **do**

2:   **for**  $j \in \{1, 2, 3\}$  **do**

3:     Sample  $\hat{\theta}(j) \sim \text{Beta}(\alpha_j, \beta_j)$

4:   **end for**

$$l(t) \leftarrow \arg \max_{j \in \{1, 2, 3\}} \hat{\theta}(j)$$

$$\alpha_{l(t)} \leftarrow \alpha_{l(t)} + r_{l(t)}$$

$$\beta_{l(t)} \leftarrow \beta_{l(t)} + 1 - r_{l(t)}$$

5: **end for**

---

# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# Setting

- Suppose we obtain the Bernoulli distribution parameters from an oracle, which are shown in the following table below.

Arm $j$	1	2	3
$\theta_j$	0.7	0.5	0.4

- Choose  $N = 5000$  and compute the theoretically maximized expectation of aggregate rewards over  $N$  time slots. We call it the oracle value. Note that these parameters  $\theta_j, j = 1, 2, 3$  and oracle values are unknown to the three bandit algorithms.

# Setting

Implement three bandit algorithms with following settings:

- $\epsilon$ -greedy with  $\epsilon = 0.1, 0.5, 0.9$ .
- UCB with  $c = 1, 5, 10$ .
- Thompson Sampling with
  - $\{(\alpha_1, \beta_1) = (1, 1), (\alpha_2, \beta_2) = (1, 1), (\alpha_3, \beta_3) = (1, 1)\}$  and
  - $\{(\alpha_1, \beta_1) = (601, 401), (\alpha_2, \beta_2) = (401, 601), (\alpha_3, \beta_3) = (2, 3)\}$

# Tasks

- Each experiment lasts for  $N = 5000$  time slots, and we run each experiment 200 trials. Results are averaged over these 200 independent trials.
- Compute the gaps between the algorithm outputs (aggregated rewards over  $N$  time slots) and the oracle value. Compare the numerical results of  $\epsilon$ -greedy, UCB, and Thompson Sampling. Which one is the best? Then discuss the impacts of  $\epsilon$ ,  $c$ , and  $\alpha_j$ ,  $\beta_j$  respectively.
- Give your understanding of the exploration-exploitation trade-off in bandit algorithms.
- We implicitly assume the reward distribution of three arms are independent. How about the dependent case? Can you design an algorithm to exploit such information to obtain a better result?

# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# Bayesian Bandits

- There are two arms which may be pulled repeatedly in any order.
- Each pull may result in either a success or a failure.
- The sequence of successes and failures which results from pulling arm  $i$  ( $i \in \{1, 2\}$ ) forms a Bernoulli process with unknown success probability  $\theta_i$ .
- A success at the  $t^{th}$  pull yields a reward  $\gamma^{t-1}$  ( $0 < \gamma < 1$ ), while an unsuccessful pull yields a zero reward.
- At time zero, each  $\theta_i$  has a Beta prior distribution with two parameters  $\alpha_i, \beta_i$  and these distributions are independent for different arms.

# Bayesian Bandits

- These prior distributions are updated to posterior distributions as arms are pulled.
- Since the class of Beta distributions is closed under Bernoulli sampling, posterior distributions are all Beta distributions.
- How should the arm to pull next in each time slot be chosen to maximize the total expected reward from an infinite sequence of pulls?

# Tasks

- One intuitive policy suggests that in each time slot we should pull the arm for which the current expected value of  $\theta_i$  is the largest. This policy behaves very good in most cases. Please design simulations to check the behavior of this policy.
- However, such intuitive policy is unfortunately not optimal. Please provide an example to show why such policy is not optimal.

# Tasks

- For the expected total reward under an optimal policy, show that the following recurrence equation holds:

$$R_1(\alpha_1, \beta_1) = \frac{\alpha_1}{\alpha_1 + \beta_1} [1 + \gamma R(\alpha_1 + 1, \beta_1, \alpha_2, \beta_2)]$$

$$+ \frac{\beta_1}{\alpha_1 + \beta_1} [\gamma R(\alpha_1, \beta_1 + 1, \alpha_2, \beta_2)];$$

$$R_2(\alpha_2, \beta_2) = \frac{\alpha_2}{\alpha_2 + \beta_2} [1 + \gamma R(\alpha_1, \beta_1, \alpha_2 + 1, \beta_2)]$$

$$+ \frac{\beta_2}{\alpha_2 + \beta_2} [\gamma R(\alpha_1, \beta_1, \alpha_2, \beta_2 + 1)];$$

$$R(\alpha_1, \beta_1, \alpha_2, \beta_2) = \max \{R_1(\alpha_1, \beta_1), R_2(\alpha_2, \beta_2)\}.$$

- For the above equations, how to solve it exactly or approximately?
- Find the optimal policy.

# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# Core of AlphaGo & AlphaZero

- 2016-2018: Google DeepMind team developed AlphaGo & AlphaZero System, where the core algorithm is MCTS(Monte Carlo Tree Search) + Neural Network
- Core of MCTS: UCT (multi-stage UCB)



*After humanity spent thousands of years improving our tactics, computers tell us that humans are completely wrong... I would go as far as to say not a single human has touched the edge of the truth of Go.*

Ke Jie,  
9 dan Go player



*robots will never understand the beauty of the game the same way that we humans do*

Lee Sedol,  
9 dan Go player



# Outline

- 1 Bandit Background
- 2 Multi-Armed Bandit Algorithms
- 3 Part I: Basic Bandits (Required)
- 4 Part II: Bayesian Bandits (Optional)
- 5 Part III: Monte Carlo Tree Search (Optional)
- 6 References

# References

- Reinforcement Learning: An Introduction (second edition), R. Sutton & A. Barto, 2018.
- Bandit Algorithms, T. Lattimore & C. Szepesvari, 2020.

# Remark

